



Screening rate by population, and HK containment within source support



Aggregate screening prevalence was nearly unchanged between the clean-sensor and March-October populations, but this is NOT a matched counterfactual: the periods differ in time of year and operating conditions, not only sensor integrity, and identical aggregate rates may conceal different row-level predictions.

The Saola figure rests on 475 of 662 eligible rows with 187 excluded (186 missing irr_meas, 1 missing power_W), so the retained subset may be selective; it is evidence neither for nor against PV-event detection.

Most HK rows lie within broad source support, so gross out-of-range extrapolation alone is unlikely to explain the very high screening rate; local and conditional distribution shift remain possible and are not excluded, and approximately 10-14% fall outside occupied training cells.